

Pushpak Vootla

Azure Data Engineer

Chennai, Tamil Nadu, India
vootlapushpakwork2001@gmail.com
pushpakvootla.cloud

PROFESSIONAL SUMMARY

Azure Data Engineer with 3+ years of professional experience and hands-on implementation experience across Azure Data Factory, Azure Databricks, PySpark, Delta Lake, ADLS Gen2, Unity Catalog and Microsoft Fabric.

Focused on metadata-driven orchestration, Landing-to-Raw processing, validation, framework logging, watermark controls, rerun behaviour and reliable data delivery for downstream analytics and Data Science workloads.

CORE TECHNICAL SKILLS

Data Orchestration: Azure Data Factory, Fabric Data Factory, Metadata-driven pipelines, Master-child pipelines, Parameterised pipelines, Parameterised datasets

Lakehouse and Storage: ADLS Gen2, Delta Lake, Fabric Lakehouse, Parquet, Azure Blob Storage, Azure SQL Database

Governance and Security: Unity Catalog, Azure Key Vault, Managed Identity fundamentals, Secure linked services, Role-based access fundamentals, Secret separation from code

Data Processing: Azure Databricks, PySpark, SQL, Python, Mapping Data Flows, Databricks Notebooks

Data Quality and Reliability: Schema validation, Row-count validation, Duplicate detection, Null validation, Reference-data validation, Quarantine handling, Rejected-record handling, Audit metadata, Framework logging

PROFESSIONAL EXPERIENCE

Associate Data Engineer | Ensono

January 2023 — Present | Chennai, India

Working in an enterprise technology environment with a current focus on Azure data engineering, production-oriented pipeline implementation and operational reliability.

- Contributed to metadata-driven Azure Data Factory orchestration.
- Worked with Azure Databricks and PySpark for Landing-to-Raw processing.
- Implemented and validated Delta Lake data-processing patterns.
- Added structured notebook-result handling between Databricks and ADF.
- Worked on stage-level logging, watermark protection and rerun behaviour.
- Applied runtime, file, schema, row-count and partition validation.
- Supported Unity Catalog registration and downstream data consumption.
- Supported Data Science readiness through consistent quarter_code metadata.

SELECTED AZURE DATA ENGINEERING PROJECTS

Quarterly Policy Data Delivery and Lakehouse Processing Platform

Azure Data Factory + Azure Databricks

A completed quarterly source-to-Raw platform combining metadata-driven ADF orchestration, Databricks validation, Delta Lake processing, Unity Catalog registration and protected watermark updates.

- Metadata-driven quarterly orchestration
- Structured ADF and Databricks result integration

Technology: Azure Data Factory, Azure Databricks, PySpark, Delta Lake, Unity Catalog, ADLS Gen2

Case study: pushpakvootla.cloud/projects/quarterly-policy-data-delivery-platform

Validation-First Lakehouse Ingestion with Microsoft Fabric

Microsoft Fabric

A Microsoft Fabric ingestion workflow that discovers incoming files, validates trusted office CSV files and loads approved records into a curated Lakehouse table while routing unsupported or schema-incompatible files to quarantine.

- Uses office_*.csv as the trusted filename contract
- Validates schema before curated loading

Technology: Microsoft Fabric, Fabric Data Factory, Fabric Lakehouse, Fabric Notebook, PySpark, Delta Lake

Case study: pushpakvootla.cloud/projects/validation-first-fabric-lakehouse-ingestion

Metadata-Driven Multi-Source Ingestion with Azure Data Factory

Azure Data Factory

A configurable Azure Data Factory ingestion framework that reads source behaviour from metadata and processes SQL Server, HTTP or REST and file-based sources through reusable pipeline components.

- Reads source behaviour from metadata stored in ADLS
- Iterates through configured sources using a ForEach activity

Technology: Azure Data Factory, ADLS Gen2, SQL Server, HTTP / REST, Azure SQL Database, Azure Logic Apps

Case study: pushpakvootla.cloud/projects/metadata-driven-adf-ingestion

Retail Data Validation Pipeline with Azure Databricks

Azure Databricks

A retail data pipeline that ingests order and order-item data from Amazon S3, stages it in ADLS Gen2 and applies Databricks PySpark validation before loading approved records into Azure SQL Database.

Ingests order and order-item datasets from Amazon S3

Detects duplicate order_id values through PySpark

Technology: Azure Data Factory, Azure Databricks, PySpark, Amazon S3, ADLS Gen2, Azure SQL Database

Case study: pushpakvootla.cloud/projects/retail-databricks-validation-pipeline

Movie Analytics Transformation Pipeline with Azure Data Factory

Azure Data Factory

A parameterised Azure Data Factory pipeline that reads movie data from Azure Blob Storage, applies Mapping Data Flow transformations and routes validated outputs to Azure SQL Database or ADLS Gen2 according to release-year rules.

Extracts the primary genre using a derived-column transformation

Groups movie records by genre and release year

Technology: Azure Data Factory, Mapping Data Flow, Azure Blob Storage, ADLS Gen2, Azure SQL Database, CSV

Case study: pushpakvootla.cloud/projects/movie-analytics-adf-pipeline

EARLY-CAREER TRAINING AND PROJECT EXPERIENCE

Trainee | StackRoute

March 2022 — September 2022 | Bengaluru, India

Completed a Wipro-assigned Java Full Stack and AWS Cloud training programme with hands-on development exercises.

Project Trainee — Cloud and Application Development | Wipro

February 2022 — September 2022 | Chennai, India

Worked on an early-career application-development project involving Java, MySQL and AWS deployment services.

CERTIFICATIONS

Microsoft Certified: Azure Data Fundamentals | Microsoft

Microsoft Certified: Azure AI Fundamentals | Microsoft

AWS Certified Data Engineer — Associate | Amazon Web Services

Jira Fundamentals Badge | Atlassian

EDUCATION

Bachelor of Technology, Mechanical Engineering

Hindustan University | May 2018 — May 2022

PROFESSIONAL LINKS

Portfolio: pushpakvootla.cloud